What did we do last week?

TF-IDF

Princeton AI4ALL

Today’s Challenge

- How to find the **relevance** of a word in a document?

Relevance: Problem Setting

- Ex: You have 10,000 newspaper articles. Find the articles most relevant to some phrase (e.g. “the brown cow”).
- How could you approach this problem?
  - Discuss ideas with a partner!

Relevance: Ideas

- Return any article which contains all the keywords
- Return any article containing one keyword or more
- Return the article with the most mentions of the keywords
- Other ideas?

Relevance: Challenges

- Filter out common words like “the”
  - How to do this?
- Don’t favor longer articles
  - How to do this?

Relevance: Term Frequency (TF)

- The proportion of words which equal the keyword
- Calculate the TF of “cow” for these examples:
    - “The brown cow mooed”
    - “Cow is a cow”
    - “Cow cow cow cow cow”
- Higher is better

Relevance: Inverse Document Frequency (IDF)

- What proportion of documents actually contain this word?
    - “the” gives a high value
    - “cow” gives a low value
- Take the reciprocal of this, and take the log base 2
    - “the” gives a low value, as desired
    - “cow” gives a high value

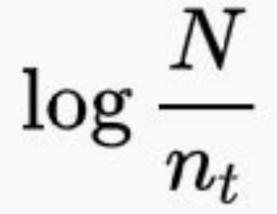

$$\log \frac{N}{n_t}$$

Relevance: TF-IDF

- Combine the two scores by multiplying TF with IDF
- Then, add the score for each word in the query
- Example 1:
  - Query: “cow”
  - Documents: “the cow moos”, “the dog barks”, “dog dog”
  - TF for “cow”: ⅓, 0, 0
  - IDF for “cow”: log(3/1)≈1.58
  - TF-IDF scores: 0.52, 0, 0
- Example 2:
  - Query: “the cow”
  - Documents: “the cow moos”, “the dog barks”, “dog dog”
  - TF for “the”: ⅓, ⅓, 0
  - IDF for “the”: log(3/2)≈0.58
  - Total TF-IDF scores: 0.19 + 0.52 = 0.71, 0 + 0.19 = 0.19, 0 + 0 = 0

Practice Time!

Practice Problems:

- Calculate the TF-IDF scores for each of the following:

Phrase: “panther”
Documents:

1) “I like that panther.”
2) “Panther panther panther.”
3) “TF-IDF is awesome.”
4) “Nothing here.”
5) “Hello.”

Phrase: “the computer broke”
Documents:

1) “We fix computers, we’re the best!”
2) “If you’re broke, you can’t buy that.”
3) “My hamster broke the computer screen.”
4) “The cat is cute.”
5) “The book is funny.”

Practice Coding

- Time to code TF-IDF in Python!

COSINE SIMILARITY

Feedback

[link to google poll]